

ICE2607 Lab 5: Pytorch & CNN

SJTU-SEIEE cny123222

December 26, 2024

1 Experiment Overview (实验概览)

Nowadays, neural networks play a crucial role in feature extraction and retrieval tasks in multimedia applications. In this experiment, we focus on utilizing neural networks for image feature extraction and retrieval, emphasizing their significance in modern multimedia retrieval systems.

The experiment is divided into two main parts. In the first part, a CIFAR-10 classifier was trained using a ResNet-20 architecture. We investigated the impact of different hyperparameters and image augmentation techniques in order to improve the classification accuracy.

In the second part, a pre-trained ResNet-50 model was employed for image retrieval on a custom dataset. Additionally, we explored the performance of the ViT-B/16 model in the same retrieval task.

2 Solution Approach (解决思路)

2.1 Part I: Image Classification

In the experiment, a ResNet-20 model was provided, requiring completion of a code snippet. Furthermore, enhancements were implemented in the codebase, aiming to facilitate a more convenient training process and improve the precision of image classification.

2.1.1 Completion of Test Accuracy Calculation

The calculation of test accuracy mirrors that of train accuracy, emphasizing the counting of accurately classified images and the total number of images. The provided code snippet illustrates this process.

```
1 model.eval()
2 correct = 0 # number of correctly classified images
3 total = 0 # total number of images
4 with torch.no_grad():
5     for batch_idx, (inputs, targets) in enumerate(testloader):
6         outputs = model(inputs)
7         _, predicted = outputs.max(1)
8         total += targets.size(0)
9         correct += predicted.eq(targets).sum().item()
10 test_accuracy = 100 * correct / total
```

```
11 print(f"Epoch [{epoch}] - Test Accuracy: {test_accuracy:.3f}%")
```

2.1.2 Improvements to Training Process

- **Utilizing AdamW over SGD:** As a more commonly used optimizer, AdamW offers greater stability in parameter updates compared to SGD, effectively handling sparse gradients and non-stationary objectives, thus accelerating model convergence.
- **Incorporating a Scheduler:** The implementation of a scheduler allows for convenient control over learning rate variations.
- **Leveraging wandb for Training Monitoring:** Employing wandb for tracking the training process not only aids in monitoring model training but also facilitates automatic hyperparameter optimization. This real-time monitoring and parameter tuning can expedite the identification of the optimal model configuration.
- **Utilizing GPU Acceleration:** Training on GPUs significantly boosts training speed, enabling us to train for more epochs and swiftly experiment with various hyperparameter combinations.

2.2 Part II: Image Retrieval

Since the feature extraction process has been provided, our experiment mainly focused on the implementation of similarity computation and a top-5 sorting mechanism.

2.2.1 Similarity Computation

When provided with a vector representing the feature of an input image and a matrix where each row portrays the feature of a dataset image, both cosine similarity and Euclidean similarity can be easily implemented.

```
1 def cosine_similarity(input_feature, feature_vectors):
2     input_feature = input_feature / np.linalg.norm(input_feature) # input_feature.shape =
                               (D,)
3     feature_vectors = feature_vectors / np.linalg.norm(feature_vectors, axis=1,
                               keepdims=True) # feature_vectors.shape = (N, D)
4     similarities = np.dot(feature_vectors, input_feature.T) # similarities.shape = (D, 1)
5     return similarities.flatten()
6
7 def euclidean_similarity(input_feature, feature_vectors):
8     input_feature = input_feature / np.linalg.norm(input_feature)
9     feature_vectors = feature_vectors / np.linalg.norm(feature_vectors, axis=1,
                               keepdims=True)
10    distances = np.linalg.norm(feature_vectors - input_feature, axis=1)
11    similarities = 1 / (1 + distances) # Adjust similarities to [0, 1]
12    return similarities
```

2.2.2 Top5 Sorting

When the dataset consists of a small number of images, sorting to identify the top 5 similarities is straightforward.

```
1 # similarities contains similarity values between input image and all dataset images
2 top5_indices = np.argsort(similarities)[::-1][:5]
3 top5_similarity_values = similarities[top5_indices]
4 top5_image_paths = image_paths[top5_indices]
```

However, in the case of a large dataset, sorting the similarity vectors can be time-consuming. In such scenarios, a priority queue can be employed for efficient sorting.

3 Experimental Results (实验结果)

3.1 Experimental Environment

The experiments were conducted on macOS Sonoma 14.6.1 with Python 3.12 and an Apple M2 Pro GPU, utilizing PyTorch 2.5.1. Additionally, a portion of the experiments was run on a server equipped with PyTorch 2.5.1, Python 3.12 (Ubuntu 22.04), CUDA 12.4, and an RTX 4090D GPU.

3.2 Classification Results

The ResNet-20 model was trained using a learning rate of 0.01, a batch size of 128, the AdamW optimizer with a weight decay of 0.01, and a StepLR scheduler with a step size of 40 and a gamma of 0.1. As Figure 1 and Figure 2 shows, after 50 epochs of training, the loss decreased to 0.108, the training accuracy improved to 96.234%, and the test accuracy reached **90.930%**.




Figure 1 Loss Curve




Figure 2 Training and Test Accuracy

3.3 Retrieval Results

By randomly selecting images from the 10 classes of CIFAR-10, we created a custom dataset comprising 994 images. Figure 3 showcases a selection of example images from this dataset.




Figure 3 Example Images from the Custom Dataset

Some images not present in the custom dataset were used for retrieval purposes. The retrieval outcomes, displayed in Figure 4, exhibit a high degree of success. This is evident from the results, which effectively locate images containing similar objects.

These results highlight the capability of ResNet-50 in extracting features and emphasize the accuracy of our similarity computation methods in identifying images with similar content.

4 Analysis and Discussion (分析与思考)

4.1 Influence of Hyperparameters

Hyperparameters play a crucial role in training neural networks. In our experiment, the Sweep function within Weights & Biases was employed to identify the optimal combination of hyperparameters, focusing primarily on the selection of **learning rate and weight decay**.

4.1.1 Influence of Learning Rate

Figure 5 depicts the loss curves corresponding to different learning rates (0.1, 0.01, 0.001, 0.0001) under the AdamW optimizer, with 50 epochs, a batch size of 128, and a StepLR scheduler with a step size of 20 and a gamma of 0.1.

The graph illustrates that excessively small learning rates, such as 0.001 or 0.0001, result in slow convergence, with the training loss decreasing at a sluggish pace. Conversely, when the learning rate is too large, for instance, 0.1, overly large step sizes can cause the optimization process to overshoot and potentially miss the local minima of the loss function. Through the experiment, it was observed that **a learning rate of 0.01 strikes a balance**, facilitating efficient convergence and avoiding the pitfalls associated with rates that are either too small or too large.

4.1.2 Tuning Hyperparameters with Sweep

Utilizing the Sweep functionality, we delved into the process of fine-tuning hyperparameters to enhance the performance of our neural network model. Figure 6 presents a detailed overview of the test accuracy




Figure 4 Retrieval Results



Figure 5 Loss Curves under Different Learning Rates

outcomes obtained through the exploration of varying learning rates (0.1, 0.01, 0.001, 0.0001) and weight decays (0.01, 0.001, 0.0001) using the Sweep feature.




Figure 6 Test Accuracy with Different Learning Rates and Weight Decays

Notably, the experimental results highlight that a **learning rate of 0.01 coupled with a weight decay of 0.01** stands out as the optimal configuration, yielding the highest test accuracy among the evaluated combinations. This meticulous tuning process exemplifies the critical role that hyperparameter selection plays in maximizing the effectiveness of neural network training and ultimately enhancing model performance.

4.2 Problem of Overfitting

In Figure 2, a noticeable disparity is evident between the train accuracy and test accuracy, indicative of a prevalent issue in machine learning termed **overfitting**. Overfitting occurs when a model captures noise from the training data rather than learning the underlying patterns, leading to poor generalization on unseen data.

To mitigate overfitting, various regularization techniques such as weight decay, dropout, and data augmentation can be employed. In our experiment, **data augmentation** was implemented as a strategy to alleviate overfitting. Data augmentation involves artificially increasing the diversity of the training dataset by applying transformations such as cropping, flipping, color jittering, rotation, etc., to introduce variability

and enhance the model's ability to generalize. The modified code snippet incorporating data augmentation is presented below:

```
1 transform_train = transforms.Compose([
2     transforms.RandomCrop(32, padding=4),
3     transforms.RandomHorizontalFlip(),
4     transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1), # Add
5     color jitter
6     transforms.RandomRotation(15), # Random rotation
7     transforms.ToTensor(),
8     transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
9 ])
```

Figure 7 displays the train and test accuracy when employing data augmentation. A comparative analysis with Figure 2 underscores the effectiveness of data augmentation in minimizing the disparity between train accuracy and test accuracy, thereby notably mitigating the issue of overfitting.




Figure 7 Train and Test Accuracy with Data Augmentation

4.3 Learning Rate Adjustment

During our training process, a StepLR scheduler was utilized, so that the learning rate decayed from 0.01 to 0.001 at the 40th epoch. Upon analyzing the trends illustrated in Figures 1 and 2, it becomes evident that initially, both the loss function and accuracy metrics were converging. However, with the subsequent decay in the learning rate, the loss exhibited a further decline while the accuracy demonstrated a notable uptick.

This observed enhancement can be attributed to the strategic adjustment of the learning rate. By lowering the learning rate at a specific epoch, the model was able to navigate the optimization landscape more effectively, fine-tuning its parameters with greater precision. This adjustment likely enabled the model to overcome local minima and plateaus that might have hindered its progress, ultimately leading to improved convergence and higher accuracy on the test data.

4.4 ViT-B/16 Comparison in Image Retrieval

The vision transformer architecture has gained prominence in recent years, excelling in various tasks including image classification and image retrieval. In our experiment, we employed ViT-B/16 to undertake the same image retrieval task as ResNet-50, enabling a comparative analysis.

Figure 8 illustrates the retrieval outcomes of both ResNet-50 and ViT-B/16. Notably, ViT-B/16 outperformed ResNet-50 in this task. The images identified by ResNet-50 as having the top 5 similarities predominantly share a red color theme, even including a red truck. Conversely, the images retrieved by ViT-B/16 all feature automobiles, with some not exhibiting a red coloration. This contrast suggests that ViT-B/16 captures more high-level features beyond color attributes.



Figure 8 Comparison of Retrieval Results between ResNet-50 (Top) and ViT-B/16 (Bottom)

5 Reflection and Conclusion (实验感想)

5.1 Learning from Experiments

Throughout these experiments, our understanding of neural networks and image retrieval has been significantly enriched. We have refined our ability to select optimal hyperparameters and further polished our proficiency in LaTeX. These experiences have not only enhanced our research capabilities but have also proven invaluable for individuals venturing into scientific exploration.

5.2 Challenges Encountered

During the course of our experiments, we encountered challenges in the **selection of hyperparameters**, which had a substantial impact on the accuracy of our models. To address this issue, we leveraged the WandB sweep functionality to conduct experiments efficiently. However, we found that this process was time-consuming. In our quest for optimal hyperparameters, we even delved into distributed computing, enabling us to run multiple networks simultaneously. This exploration allowed us to expedite the parameter tuning process.

5.3 Conclusion

In this experiment, we utilized ResNet for image classification and image retrieval tasks, yielding commendable results. Our investigation focused on strategies for hyperparameter selection, techniques to combat overfitting, and the importance of learning rate adjustments. Additionally, we conducted a comparative analysis of the retrieval performance between ResNet-50 and ViT-B/16 models.

A Source Code File List

Table 1 File List

File Name	Description
train.py	Image Classification
retrieval.py	Image Retrieval

B Source Code

```
1  """
2  File name: train.py
3  """
4  import os
5  import sys
6  import argparse
7  import torch
8  import torch.nn as nn
9  import torch.optim as optim
10 import torchvision
11 import torchvision.transforms as transforms
12 import wandb
13 from model import resnet20
14
15
16 def parse_args(args):
17     """
18     Parse command line arguments
19     """
20     parser = argparse.ArgumentParser()
21     parser.add_argument("--lr", type=float, default=0.01, help="Learning rate")
22     parser.add_argument("--batch_size", type=int, default=128, help="Batch size")
23     parser.add_argument("--epochs", type=int, default=50, help="Number of training
24         epochs")
25     parser.add_argument("--optimizer", type=str, default="AdamW", choices=["SGD", "Adam",
26         "AdamW"], help="Optimizer type")
27     parser.add_argument("--weight_decay", type=float, default=0.01, help="Weight decay")
28     parser.add_argument("--step_size", type=int, default=40, help="Step size for learning
29         rate scheduler")
30     parser.add_argument("--augment", type=bool, default=False, help="Augment data")
31
32     args = parser.parse_args(args)
```

```

30     return args
31
32
33 def main(args):
34     # Get arguments
35     args = parse_args(args)
36
37     # Initialize Weights & Biases
38     wandb.init(
39         project="Lab5-ResNet20",
40         name=f"resnet20_lr_{args.lr}_batch_{args.batch_size}_epochs_{args.epochs}_optimizer_{args.optimizer}_w
41         config={
42             "lr": args.lr,
43             "batch_size": args.batch_size,
44             "epochs": args.epochs,
45             "optimizer": args.optimizer,
46             "weight_decay": args.weight_decay,
47             "step_size": args.step_size,
48             "augment": args.augment,
49         }
50     )
51     config = wandb.config
52
53     # Data pre-processing
54     print('==> Preparing data..')
55     if config.augment: # Augment data
56         transform_train = transforms.Compose([
57             transforms.RandomCrop(32, padding=4),
58             transforms.RandomHorizontalFlip(),
59             transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
60             # Add color jitter
61             transforms.RandomRotation(15), # Random rotation
62             transforms.ToTensor(),
63             transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
64         ])
65     else:
66         transform_train = transforms.Compose([
67             transforms.RandomCrop(32, padding=4),
68             transforms.RandomHorizontalFlip(),
69             transforms.ToTensor(),
70             transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
71         ])
72     transform_test = transforms.Compose([
73         transforms.ToTensor(),

```

```

73     transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
74 ]
75
76 # Get training data
77 trainset = torchvision.datasets.CIFAR10(
78     root='./data', train=True, download=True, transform=transform_train)
79 trainloader = torch.utils.data.DataLoader(trainset, batch_size=config.batch_size,
80     shuffle=True)
81
82 # Get testing data
83 testset = torchvision.datasets.CIFAR10(
84     root='./data', train=False, download=True, transform=transform_test)
85 testloader = torch.utils.data.DataLoader(testset, batch_size=config.batch_size,
86     shuffle=False)
87
88 classes = ("airplane", "automobile", "bird", "cat",
89     "deer", "dog", "frog", "horse", "ship", "truck")
90
91 # Model
92 print('==> Building model..')
93 model = resnet20()
94
95 # Use GPU if available
96 if torch.cuda.is_available():
97     device = torch.device("cuda")
98 elif torch.mps.is_available():
99     device = torch.device("mps")
100 else:
101     device = torch.device("cpu")
102 model = model.to(device)
103
104 # Loss function
105 criterion = nn.CrossEntropyLoss()
106
107 # Optimizer (AdamW by default)
108 if config.optimizer == "SGD":
109     optimizer = optim.SGD(model.parameters(), lr=config.lr,
110         weight_decay=config.weight_decay)
111 elif config.optimizer == "Adam":
112     optimizer = optim.Adam(model.parameters(), lr=config.lr,
113         weight_decay=config.weight_decay)
114 elif config.optimizer == "AdamW":
115     optimizer = optim.AdamW(model.parameters(), lr=config.lr,
116         weight_decay=config.weight_decay)

```

```

112
113 # Learning rate scheduler (StepLR)
114 scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=config.step_size,
115                                             gamma=0.1)
116
117 # Training for one epoch
118 def train(epoch):
119     model.train()
120     train_loss = 0
121     correct = 0
122     total = 0
123     for batch_idx, (inputs, targets) in enumerate(trainloader):
124         inputs, targets = inputs.to(device), targets.to(device)
125         optimizer.zero_grad()
126
127         # Forward pass
128         outputs = model(inputs)
129         loss = criterion(outputs, targets)
130         loss.backward()
131         optimizer.step()
132
133         # Calculate training accuracy
134         train_loss += loss.item()
135         _, predicted = outputs.max(1)
136         total += targets.size(0)
137         correct += predicted.eq(targets).sum().item()
138         print('Epoch [%d] Batch [%d/%d] Loss: %.3f | Training Acc: %.3f%% (%d/%d)'
139               % (epoch, batch_idx + 1, len(trainloader), train_loss / (batch_idx + 1),
140                  100. * correct / total, correct, total))
141
142     avg_train_loss = train_loss / len(trainloader)
143     train_accuracy = 100. * correct / total
144
145     return avg_train_loss, train_accuracy
146
147 # Testing for one epoch
148 def test(epoch):
149     print('==> Testing...')
150     model.eval()
151     correct = 0
152     total = 0
153     with torch.no_grad():
154         for batch_idx, (inputs, targets) in enumerate(testloader):
155             inputs, targets = inputs.to(device), targets.to(device)

```

```

155         outputs = model(inputs)
156
157         # Calculate testing accuracy
158         _, predicted = outputs.max(1)
159         total += targets.size(0)
160         correct += predicted.eq(targets).sum().item()
161
162     test_accuracy = 100 * correct / total
163     print(f"Epoch [{epoch}] - Test Accuracy: {test_accuracy:.3f}%")
164
165     # Save checkpoint
166     print('Saving..')
167     state = {
168         'net': model.state_dict(),
169         'acc': test_accuracy,
170         'epoch': epoch
171     }
172     if not os.path.isdir('checkpoint'):
173         os.mkdir('checkpoint')
174     torch.save(state, f'./checkpoint/ckpt_{epoch}_acc_{test_accuracy:.3f}.pth')
175
176     return test_accuracy
177
178
179 # Training and Testing loop
180 for epoch in range(1, config.epochs + 1):
181     avg_train_loss, train_accuracy = train(epoch)
182     test_accuracy = test(epoch)
183
184     wandb.log({
185         'train_loss': avg_train_loss,
186         'train_accuracy': train_accuracy,
187         'test_accuracy': test_accuracy,
188     })
189
190     scheduler.step()
191
192 wandb.finish()
193
194
195 if __name__ == '__main__':
196     main(sys.argv[1:])

```

```
1 """
```

```

2 File name: retrieval.py
3 """
4 import os
5 import sys
6 import glob
7 import argparse
8 import numpy as np
9 import matplotlib.pyplot as plt
10 from tqdm import tqdm
11 import torch
12 import torch.nn as nn
13 import torchvision.transforms as transforms
14 import torchvision.models as models
15 from torchvision.models import ResNet50_Weights
16 from torchvision.models import ViT_B_16_Weights
17 from torchvision.datasets import ImageFolder
18 from torch.utils.data import DataLoader
19 from PIL import Image
20
21 # Plot settings
22 plt.rcParams['font.family'] = 'serif'
23 plt.rcParams['font.size'] = 11
24 plt.rcParams['axes.labelsize'] = 11
25 plt.rcParams['xtick.labelsize'] = 11
26
27
28 def parse_args(args):
29     """
30     Parse command line arguments
31     """
32     parser = argparse.ArgumentParser()
33     parser.add_argument("--model", type=str, default="resnet50", choices=["resnet50",
34         "vit_b_16"], help="Model name")
35     parser.add_argument("--dataset-dir", type=str, default="./dataset", help="Dataset
36         directory")
37     parser.add_argument("--target-dir", type=str, default="./query", help="Target image
38         path or directory")
39     parser.add_argument("--output-dir", type=str, default="./figures/retrieval",
40         help="Output directory")
41
42     args = parser.parse_args(args)
43     return args

```

```

42 def get_image_paths(input_dir, extensions = ("jpg", "jpeg", "png", "bmp")):
43     """
44     Get image paths from the given directory
45     """
46     pattern = f"{input_dir}/**/*"
47     img_paths = []
48     for extension in extensions:
49         img_paths.extend(glob.glob(f"{pattern}.{extension}", recursive=True))
50
51     if not img_paths:
52         raise FileNotFoundError(f"No images found in {input_dir}. Supported formats are:
53             {'', '.join(extensions)}")
54
55     return img_paths
56
57 def extract_features(model, dataset_dir, trans, device):
58     """
59     Extract features from the dataset
60     """
61     print('==> Preparing image data..')
62     dataset = ImageFolder(dataset_dir, transform=trans)
63     dataloader = DataLoader(dataset, batch_size=1, shuffle=False)
64
65     print("==> Extracting features..")
66     feature_vectors = []
67     image_paths = []
68     with torch.no_grad():
69         for inputs, labels in tqdm(dataloader, desc="Extracting features"):
70             inputs = inputs.to(device)
71             features = model(inputs)
72             feature_vectors.append(features.cpu().numpy())
73             image_paths.append(dataloader.dataset.samples[len(feature_vectors) - 1][0])
74
75     feature_vectors = np.vstack(feature_vectors)
76     image_paths = np.array(image_paths)
77     return feature_vectors, image_paths
78
79
80 def cosine_similarity(input_feature, feature_vectors):
81     """
82     Calculate cosine similarity between the input feature and the feature vectors
83     """
84     input_feature = input_feature / np.linalg.norm(input_feature)

```



```

85     feature_vectors = feature_vectors / np.linalg.norm(feature_vectors, axis=1,
86         keepdims=True)
87     similarities = np.dot(feature_vectors, input_feature.T)
88     return similarities.flatten()
89
90 def euclidean_similarity(input_feature, feature_vectors):
91     """
92     Calculate similarity based on Euclidean distance between the input feature and the
93     feature vectors
94     """
95     input_feature = input_feature / np.linalg.norm(input_feature)
96     feature_vectors = feature_vectors / np.linalg.norm(feature_vectors, axis=1,
97         keepdims=True)
98     distances = np.linalg.norm(feature_vectors - input_feature, axis=1)
99     similarities = 1 / (1 + distances)
100     return similarities
101
102 def main(args):
103     args = parse_args(args)
104
105     # Load model
106     print("==> Loading model..")
107     print(f"Model: {args.model}")
108     if args.model == "resnet50":
109         model = models.resnet50(weights=ResNet50_Weights.IMAGENET1K_V1)
110     elif args.model == "vit_b_16":
111         model = models.vit_b_16(weights=ViT_B_16_Weights.IMAGENET1K_V1)
112     model.fc = nn.Identity() # Remove the classification head
113     model.eval()
114
115     # Use GPU if available
116     if torch.cuda.is_available():
117         device = torch.device("cuda")
118     elif torch.mps.is_available():
119         device = torch.device("mps")
120     else:
121         device = torch.device("cpu")
122     model = model.to(device)
123
124     # Data pre-processing
125     normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
126         std=[0.229, 0.224, 0.225])

```

```

126     trans = transforms.Compose([
127         transforms.Resize(256),
128         transforms.CenterCrop(224),
129         transforms.ToTensor(),
130         normalize,
131     ])
132
133     # Extract features
134     feature_path = f"{args.dataset_dir}/features_{args.model}.npz"
135     image_paths_path = f"{args.dataset_dir}/image_paths_{args.model}.npz"
136     if not os.path.exists(feature_path) or not os.path.exists(image_paths_path):
137         features, image_paths = extract_features(model, args.dataset_dir, trans, device)
138         print('==> Saving features..')
139         np.savez(feature_path, features)
140         np.savez(image_paths_path, image_paths)
141     else:
142         print('==> Loading features..')
143         features = np.load(feature_path)
144         image_paths = np.load(image_paths_path)
145
146     # Get target image paths
147     if os.path.isfile(args.target_dir):
148         target_paths = [args.target_dir]
149     else:
150         target_paths = get_image_paths(args.target_dir)
151
152     for target_path in target_paths:
153         print()
154         print(f"Query Image: {target_path}")
155
156         # Preprocess the target image
157         target_img = Image.open(target_path).convert('RGB')
158         target_img = trans(target_img).unsqueeze(0)
159         target_img = target_img.to(device)
160
161         # Extract features from the target image
162         with torch.no_grad():
163             input_feature = model(target_img).cpu().numpy()
164
165         # Calculate cosine similarity
166         similarities = cosine_similarity(input_feature, features)
167         # similarities = euclidean_similarity(input_feature, features)
168
169         # Get top-5 similar images

```

```

170     top5_indices = np.argsort(similarities)[:,-1][:5]
171     print("Top-5 Similar Indices:", top5_indices)
172     top5_similarities = similarities[top5_indices]
173     top5_image_paths = image_paths[top5_indices]
174     print("Top-5 Similar Images:")
175     for i, (path, sim) in enumerate(zip(top5_image_paths, top5_similarities)):
176         print(f"{i + 1}: {path}, Similarity: {sim:.4f}")
177
178     # Visualize the results
179     plt.figure(figsize=(10, 2.5))
180     plt.subplot(1, 6, 1)
181     query_img = plt.imread(target_path)
182     plt.imshow(query_img)
183     plt.title("Query Image\n" + target_path, fontsize=9, color='red')
184     plt.axis("off")
185     for i, path in enumerate(top5_image_paths):
186         plt.subplot(1, 6, i + 2)
187         img = plt.imread(path)
188         plt.imshow(img)
189         plt.title(f"Rank {i + 1}\nSimilarity: {top5_similarities[i]:.4f}\n{path}",
190                 fontsize=9)
191         plt.axis("off")
192     plt.tight_layout()
193     save_path =
194         f"{args.output_dir}/retrieval_{args.model}_{os.path.basename(target_path)}"
195     plt.savefig(save_path, dpi=300)
196     print("Results saved to:", save_path)
197
198 if __name__ == '__main__':
199     main(sys.argv[1:])

```